







LLM Interview Questions



Crack AI & GenAI Roles with Confidence



Naresh Edagotti
Follow For More

What is tokenization, and why is it important in LLMs?

ANSWER

Tokenization is the process of splitting text into smaller units called tokens, which can be words, subwords, or even characters. For instance, the word 'tokenization' might be broken down into smaller subwords like 'token' and 'ization'. This step is crucial because LLMs do not understand raw text directly. Instead, they process sequences of numbers that represent these tokens. Effective tokenization allows models to handle various languages, manage rare words, and reduce the vocabulary size, which improves both efficiency and performance.

CHECK **CAPTION** FOR **AI INTERVIEW** BUNDLE

What is LoRA and QLoRA?

ANSWER

LoRA and QLoRA are techniques designed to optimize the fine-tuning of Large Language Models (LLMs), focusing on reducing memory usage and enhancing efficiency without compromising performance in Natural Language Processing (NLP) tasks. LoRA (Low-Rank Adaptation) is a parameter-efficient fine-tuning method that introduces new trainable parameters to modify a model's behavior without increasing its overall size. QLoRA (Quantized LoRA) builds on LoRA by incorporating quantization to further optimize memory usage, using techniques such as 4-bit Normal Float, Double Quantization, and Paged Optimizers to compress the model's parameters.

What is beam search, and how does it differ from greedy decoding?

ANSWER

Beam search is a search algorithm used during text generation to find the most likely sequence of words. Instead of choosing the single highest-probability word at each step (as greedy decoding does), beam search explores multiple possible sequences in parallel, maintaining a set of the top k candidates (beams). It balances between finding high-probability sequences and exploring alternative paths. This leads to more coherent and contextually appropriate outputs, especially in long-form text generation tasks.

Explain the concept of temperature in LLM text generation.

ANSWER

Temperature is a hyperparameter that controls the randomness of text generation by adjusting the probability distribution over possible next tokens. A low temperature (close to 0) makes the model highly deterministic, favoring the most probable tokens. Conversely, a high temperature (above 1) encourages more diversity by flattening the distribution, allowing less probable tokens to be selected. For instance, a temperature of 0.7 strikes a balance between creativity and coherence, making it suitable for generating diverse but sensible outputs.

What is masked language modeling, and how does it contribute to model pretraining?

ANSWER

Masked language modeling (MLM) is a training objective where some tokens in the input are randomly masked, and the model is tasked with predicting them based on context. This forces the model to learn contextual relationships between words, enhancing its ability to understand language semantics. MLM is commonly used in models like BERT, which are pretrained using this objective to develop a deep understanding of language before fine-tuning on specific tasks.

What are Sequence-to-Sequence Models?

ANSWER

Sequence-to-Sequence (Seq2Seq) Models are a type of neural network architecture designed to transform one sequence of data into another sequence. These models are commonly used in tasks where the input and output have variable lengths, such as in machine translation, text summarization, and speech recognition.

How do autoregressive models differ from masked models in LLM training?

ANSWER

Autoregressive models, such as GPT, generate text one token at a time, with each token predicted based on the previously generated tokens. This sequential approach is ideal for tasks like text generation. Masked models, like BERT, predict randomly masked tokens within a sentence, leveraging both left and right context.

Autoregressive models excel in generative tasks, while masked models are better suited for understanding and classification tasks.

What role do embeddings play in LLMs, and how are they initialized?

ANSWER

Embeddings are dense, continuous vector representations of tokens, capturing semantic and syntactic information. They map discrete tokens (words or subwords) into a high-dimensional space, making them suitable for input into neural networks. Embeddings are typically initialized randomly or with pretrained vectors like Word2Vec or GloVe. During training, these embeddings are fine-tuned to capture task-specific nuances, enhancing the model's performance on various language tasks.

What is next sentence prediction and how is useful in language modelling?

ANSWER

Next Sentence Prediction (NSP) is a key technique used in language modeling, particularly in training large models like BERT (Bidirectional Encoder Representations from Transformers). NSP helps a model understand the relationship between two sentences, which is important for tasks like question answering, dialogue generation, and information retrieval. During pre-training, the model is fed two sentences: 50% of the time, the second sentence is the actual next sentence (positive pairs), and 50% of the time, it is a random sentence (negative pairs). The model is trained to classify whether the second sentence is correct or not, improving its overall language understanding.

Explain the difference between top-k sampling and nucleus (top-p) sampling in LLMs.

ANSWER

Top-k sampling restricts the model's choices to the top k most probable tokens at each step, introducing controlled randomness. For example, setting $k=10$ means the model will only consider the 10 most likely tokens. Nucleus sampling, or top-p sampling, takes a more dynamic approach by selecting tokens whose cumulative probability exceeds a threshold p (e.g., 0.9). This allows for flexible candidate sets based on context, promoting both diversity and coherence in generated text.

How does prompt engineering influence the output of LLMs?

ANSWER

Prompt engineering involves crafting input prompts to guide an LLM's output effectively. Since LLMs are highly sensitive to input phrasing, a well-designed prompt can significantly influence the quality and relevance of the response. For example, adding context or specific instructions within the prompt can improve accuracy in tasks like summarization or question-answering. Prompt engineering is especially useful in zero-shot and few-shot learning scenarios, where task-specific examples are minimal.

How can catastrophic forgetting be mitigated in large language models (LLMs)?

ANSWER

Catastrophic forgetting happens when an LLM forgets previously learned tasks while learning new ones, which limits its versatility. To mitigate this, several strategies are used: Rehearsal methods involve retraining the model on a mix of old and new data, helping it retain knowledge of previous tasks. Elastic Weight Consolidation (EWC) assigns importance to certain model weights, protecting critical knowledge while learning new tasks. Modular approaches like progressive neural networks (Prog Net) and optimized fixed expansion layers (OFELs) introduce new modules for new tasks, allowing the LLM to learn without overwriting prior knowledge.

What is model distillation, and how is it applied to LLMs?

ANSWER

Model distillation is a technique where a smaller, simpler model (student) is trained to replicate the behavior of a larger, more complex model (teacher). In the context of LLMs, the student model learns from the teacher's soft predictions rather than hard labels, capturing nuanced knowledge. This approach reduces computational requirements and memory usage while maintaining similar performance, making it ideal for deploying LLMs on resource-constrained devices.

How do LLMs handle out-of-vocabulary (OOV) words?

ANSWER

Out-of-vocabulary words refer to words that the model did not encounter during training. LLMs address this issue through subword tokenization techniques like Byte-Pair Encoding (BPE) and WordPiece. These methods break down OOV words into smaller, known subword units. For example, the word 'unhappiness' might be tokenized as 'un,' 'happi,' and 'ness'. This allows the model to understand and generate words it has never seen before by leveraging these subword components.

How does the Transformer architecture overcome the challenges faced by traditional Sequence-to-Sequence models?

ANSWER

The Transformer architecture overcomes key limitations of traditional Seq2Seq models in several ways: Parallelization - Transformers use self-attention to process tokens in parallel, speeding up both training and inference. Long-Range Dependencies - Transformers capture these effectively with self-attention, allowing the model to focus on any part of the sequence, regardless of distance. Positional Encoding ensures the model understands token order. Efficiency and Scalability - Transformers scale better for large datasets and long sequences. Context Bottleneck - Transformers let the decoder attend to all encoder outputs, improving context retention.

What is overfitting in machine learning, and how can it be prevented?

ANSWER

Overfitting occurs when a machine learning model performs well on the training data but poorly on unseen or test data. This typically happens because the model has learned not only the underlying patterns in the data but also the noise and outliers. Techniques to overcome overfitting include: Regularization (L1, L2) - Adding a penalty to the loss function. Dropout - Randomly deactivating neurons during training. Data Augmentation - Expanding the training dataset with variations. Early Stopping - Stopping training when validation loss stops decreasing. Simpler Models - Reducing complexity by decreasing features, parameters, or layers.

What are Generative and Discriminative models?

ANSWER

In NLP, generative and discriminative models are two key types of models used for various tasks. Generative models learn the underlying data distribution and generate new samples from it. They model the joint probability distribution of inputs and outputs. A common example is a language model predicting the next word. Discriminative models focus on learning a decision boundary between different classes. They model the conditional probability of outputs given inputs to accurately classify new examples. An example is a sentiment analysis model.

How is GPT-4 different from its predecessors like GPT-3 in terms of capabilities and applications?

ANSWER

GPT-4 introduces several advancements over GPT-3:

Improved Understanding - GPT-4 has roughly 1 trillion parameters, significantly more than GPT-3's 175 billion.

Multimodal Capabilities - GPT-4 processes both text and images, unlike the text-only GPT-3.

Larger Context Window - GPT-4 can handle up to 25,000 tokens, while GPT-3 maxes out at 4,096 tokens.

Better Accuracy and Fine-Tuning - GPT-4 is more factually accurate and produces less harmful information.

Language Support - GPT-4 has improved multilingual performance, supporting up to 26 languages with higher accuracy.

What are positional encodings in the context of large language models?

ANSWER

Positional encodings are essential in LLMs to address the inability of transformer architectures to capture sequence order. Since transformers process tokens simultaneously through self-attention, they are unaware of token order. Positional encodings provide the necessary information to help the model understand the sequence of words.

Additive Approach - Encodings are added to input word embeddings, merging static representations with positional data.

Sinusoidal Function - Models use trigonometric functions (Sine and Cosine) to generate these encodings based on position and dimension.

What is Multi-head attention?

ANSWER

Multi-head attention is an enhancement of single-head attention, allowing a model to attend to information from different representation subspaces simultaneously. Instead of using a single attention mechanism, multi-head attention projects the queries, keys, and values into multiple subspaces through distinct learned linear transformations. The attention function is applied in parallel to each projected version, generating multiple output vectors that are then combined to produce the final result. This approach improves the model's ability to capture complex patterns and relationships.

Derive the softmax function and explain its role in attention mechanisms.

ANSWER

The softmax function transforms a vector of real numbers into a probability distribution. For an input vector x , the softmax function for the i -th element is defined as: $\text{softmax}(x_i) = \exp(x_i) / \sum \exp(x_j)$. This ensures all output values lie between 0 and 1 and sum to 1, making them interpretable as probabilities. In attention mechanisms, softmax is applied to the attention scores to normalize them, allowing the model to assign varying levels of importance to different tokens when generating output. This helps the model focus on the most relevant parts of the input sequence.

How is the dot product used in self-attention, and what are its implications for computational efficiency?

ANSWER

In self-attention, the dot product is used to calculate the similarity between query (Q) and key (K) vectors. The attention scores are computed as: $\text{Attention}(Q, K, V) = \text{softmax}(QK^T / \sqrt{d_k})V$, where d_k is the dimensionality of the key vectors. The dot product measures alignment between tokens, helping the model decide which tokens to focus on. While effective, the quadratic complexity $O(n^2)$ in sequence length can be a challenge for long sequences, prompting the development of more efficient approximations.

Explain cross-entropy loss and why it is commonly used in language modeling.

ANSWER

Cross-entropy loss measures the difference between the predicted probability distribution and the true distribution (one-hot encoding of the correct token). It is defined as: $L = -\sum y_i * \log(\hat{y}_i)$, where y_i is the true label and $\hat{y}_i$ is the predicted probability. Cross-entropy loss penalizes incorrect predictions more heavily, encouraging the model to output probabilities closer to 1 for the correct class. In language modeling, it ensures the model predicts the correct token in a sequence with high confidence.

How do you compute the gradient of the loss function with respect to embeddings?

ANSWER

To compute the gradient of the loss L with respect to an embedding vector e , you apply the chain rule: $\partial L / \partial e = (\partial L / \partial \text{logits}) * (\partial \text{logits} / \partial e)$. Here, $\partial L / \partial \text{logits}$ is the gradient of the loss with respect to the output logits, and $\partial \text{logits} / \partial e$ is the gradient of the logits with respect to the embeddings. Backpropagation propagates these gradients through the network layers, adjusting the embedding vectors to minimize the loss.

What is the role of the Jacobian matrix in backpropagation through a transformer model?

ANSWER

The Jacobian matrix represents the partial derivatives of a vector-valued function with respect to its inputs. In backpropagation, it captures how each element of the output vector changes with respect to each input. For transformer models, the Jacobian is essential in computing gradients for multi-dimensional outputs, ensuring that each parameter (including weights and embeddings) is updated correctly to minimize the loss function.

Explain the concept of eigenvalues and eigenvectors in the context of matrix factorization for dimensionality reduction.

ANSWER

Eigenvalues and eigenvectors are fundamental in understanding the structure of matrices. Given a matrix A , an eigenvector v and eigenvalue λ satisfy the equation: $Av = \lambda v$. In dimensionality reduction techniques like PCA (Principal Component Analysis), eigenvectors represent the principal components, and eigenvalues indicate the amount of variance captured by each component. Selecting components with the largest eigenvalues helps reduce dimensionality while preserving most of the data's variance.

How is the KL divergence used in evaluating LLM outputs?

ANSWER

KL (Kullback-Leibler) divergence measures the difference between two probability distributions P (true distribution) and Q (predicted distribution): $KL(P||Q) = \sum P(x) * \log(P(x)/Q(x))$. In LLMs, it quantifies how much the predicted distribution deviates from the target distribution. A lower KL divergence indicates that the model's predictions closely match the true labels, making it a useful metric in evaluating and fine-tuning language models.

Derive the formula for the derivative of the ReLU activation function and discuss its significance.

ANSWER

The ReLU (Rectified Linear Unit) function is defined as: $f(x) = \max(0, x)$. Its derivative is: $f'(x) = 1$ if $x > 0$, else 0. ReLU introduces non-linearity while maintaining computational efficiency. Its sparsity (outputting zero for negative inputs) helps mitigate the vanishing gradient problem, making it a popular choice in deep learning models, including LLMs.

What is the chain rule in calculus, and how does it apply to gradient descent in deep learning?

ANSWER

The chain rule states that the derivative of a composite function $f(g(x))$ is: $d/dx[f(g(x))] = f'(g(x)) * g'(x)$. In deep learning, the chain rule is used in backpropagation to compute gradients of the loss function with respect to each parameter layer by layer. This allows gradient descent to update weights efficiently, propagating error signals backward through the network.

How do you compute the attention scores in a transformer, and what is their mathematical interpretation?

ANSWER

Attention scores in a transformer are computed as:
$$\text{Attention}(Q, K, V) = \text{softmax}(QK^T / \sqrt{d_k})V.$$
 Here, Q (queries), K (keys), and V (values) are learned representations of the input. The dot product QK^T measures the similarity between queries and keys. Scaling by $\sqrt{d_k}$ prevents excessively large values, ensuring stable gradients. The softmax function normalizes these scores, emphasizing the most relevant tokens for each query, guiding the model's focus during generation.

In what ways does Gemini's architecture optimize training efficiency and stability compared to other multimodal LLMs like GPT-4?

ANSWER

Gemini's architecture optimizes training efficiency and stability compared to multimodal models like GPT-4 in several ways:

- Unified Multimodal Design** - Gemini integrates text and image processing in a single model, improving parameter sharing and reducing complexity.
- Cross-Modality Attention** - Enhanced interactions between text and images lead to better learning and stability during training.
- Data-Efficient Pretraining** - Self-supervised and contrastive learning allow Gemini to train with less labeled data, boosting efficiency.
- Balanced Objectives** - Better synchronization of text and image losses ensures stable training and smoother convergence.

What are different types of Foundation Models?

ANSWER

Foundation models are large-scale AI models trained on vast amounts of unlabeled data using unsupervised methods. They are designed to learn general-purpose knowledge that can be applied to various tasks across domains.

Common Types:

- Language Models** - Tasks: Machine translation, text summarization, question answering. Examples: BERT, GPT-3.
- Computer Vision Models** - Tasks: Image classification, object detection, image segmentation. Examples: ResNet, VGGNet.
- Generative Models** - Tasks: Creative writing, image generation, music composition. Examples: DALL-E, Imagen.
- Multimodal Models** - Tasks: Image captioning, visual question answering. Examples: PaLM, LaMDA.

How does Parameter-Efficient Fine-Tuning (PEFT) prevent catastrophic forgetting in LLMs?

ANSWER

Parameter-Efficient Fine-Tuning (PEFT) helps prevent catastrophic forgetting in LLMs by updating only a small set of task-specific parameters, while keeping most of the model's original parameters frozen. This approach allows the model to adapt to new tasks without overwriting previously learned knowledge, ensuring it retains core capabilities while learning new information efficiently.

What are the key steps involved in the Retrieval-Augmented Generation (RAG) pipeline?

ANSWER

Key steps in the Retrieval-Augmented Generation (RAG) pipeline are: Retrieval - The query is encoded and compared with precomputed document embeddings to retrieve relevant documents. Ranking - The retrieved documents are ranked based on their relevance to the query. Generation - The top-ranked documents are used as context by the LLM to generate more informed and accurate responses. This hybrid approach enhances the model's ability to produce context-aware outputs by incorporating external knowledge during generation.

How does the Mixture of Experts (MoE) technique improve LLM scalability?

ANSWER

Mixture of Experts (MoE) improves LLM scalability by using a gating function to activate only a subset of expert models (sub-networks) for each input, rather than the entire model. This selective activation reduces computational load - Only a few experts are active per query, minimizing resource usage. Maintains high performance - The model dynamically selects the most relevant experts for each input, ensuring task complexity is handled effectively. MoE enables efficient scaling of LLMs, allowing larger models with billions of parameters while controlling computational costs.

What is Chain-of-Thought (CoT) prompting, and how does it improve complex reasoning in LLMs?

ANSWER

Chain-of-Thought (CoT) prompting helps LLMs handle complex reasoning by encouraging them to break down tasks into smaller, sequential steps. This improves their performance by:

- Simulating human-like reasoning - CoT prompts the model to approach problems step-by-step, similar to how humans solve complex issues.
- Enhancing multi-step task performance - It's particularly effective for tasks involving logical reasoning or multi-step calculations.
- Increasing accuracy - By guiding the model through a structured thought process, CoT reduces errors and improves performance on intricate queries.

CoT improves LLMs' interpretability and reliability in tasks that require deeper reasoning and decision-making.

What is the difference between discriminative AI and Generative AI?

ANSWER

Predictive/Discriminative AI focuses on predicting or classifying data based on existing data. It models the conditional probability $P(y|x)$, where y is the target variable and x represents the input features. Examples include tasks like classification (e.g., image recognition), regression, and applications such as spam detection.

Generative AI focuses on generating new data samples that resemble the training data. It models the joint probability $P(x,y)$, allowing it to create new instances of data. Examples include generating text, images, music, and other content using techniques like GANs, VAEs, and GPT.

How does knowledge graph integration enhance LLMs?

ANSWER

Integrating knowledge graphs with LLMs enhances performance by adding structured, factual knowledge. Key benefits include:

- Factual accuracy** - The model can cross-check information against the knowledge graph, reducing hallucinations and improving correctness.
- Enhanced reasoning** - Knowledge graphs support logical reasoning by leveraging relationships between entities, enabling better handling of complex queries.
- Contextual understanding** - Structured data helps the model understand context and relationships, improving response quality.

This integration is particularly valuable in tasks like question answering, entity recognition, and recommendation systems.

What is zero-shot learning, and how does it apply to LLMs?

ANSWER

Zero-shot learning enables LLMs to perform tasks they haven't been explicitly trained for by leveraging their broad understanding of language and general concepts. Instead of needing task-specific fine-tuning, the model can generate relevant outputs based on the instructions provided in the prompt. For example: Text classification - The model can categorize text without specific training, simply by understanding the prompt's context. Translation or summarization - LLMs can translate or summarize text using provided instructions, even without task-specific fine-tuning. This shows the LLMs' ability to generalize across tasks, making them versatile for various applications.

How does Adaptive Softmax speed up large language models?

ANSWER

Adaptive Softmax accelerates LLMs by categorizing words into frequency groups, allowing for fewer computations for infrequent words. This approach lowers overall computational costs while preserving accuracy, making it effective for efficiently managing large vocabularies.

What is the vanishing gradient problem, and how does the Transformer architecture address it?

ANSWER

The vanishing gradient problem occurs when gradients diminish during backpropagation, preventing deep networks from learning effectively, especially in models like RNNs that handle long sequences. Transformers address this by using:

- Self-Attention Mechanism** - Captures relationships between all tokens in a sequence simultaneously, avoiding sequential dependencies and preventing gradient shrinkage over time.
- Residual Connections** - Skip connections between layers allow gradients to flow directly, ensuring they remain strong throughout backpropagation.
- Layer Normalization** - Normalizes inputs within each layer, stabilizing gradient updates and preventing vanishing or exploding gradients.

These innovations enable deep models to learn efficiently, even for long sequences, overcoming the limitations of earlier architectures.

Explain the concept of 'few-shot learning' in LLMs and its advantages.

ANSWER

Few-shot learning in LLMs is the ability of the model to understand and tackle new tasks with just a few examples. This is made possible by the model's extensive pre-training, which allows it to generalize from limited data. The main benefits include:

- Reduced Data Needs** - It requires fewer examples to perform well, minimizing the need for large, task-specific datasets.
- Increased Flexibility** - The model can easily adapt to various tasks with minimal additional training.
- Cost Efficiency** - With less need for extensive data and reduced training times, it lowers the costs associated with data collection and computational resources.

You're working on an LLM, and it starts generating offensive or factually incorrect outputs. How would you diagnose and address this issue?

ANSWER

If an LLM produces offensive or inaccurate outputs, I would first analyze the patterns, check input prompts, and assess if the issue stems from biases or gaps in the training data. I'd review the preprocessing pipeline for errors or biases and examine the dataset for imbalances. Next, I'd evaluate the model's architecture, hyperparameters, and fine-tuning to identify any structural issues. Solutions could include adversarial training, debiasing, data augmentation, or retraining with a more balanced dataset.

How is the encoder different from the decoder?

ANSWER

In the transformer architecture used in large language models, the encoder and decoder serve different purposes. The encoder processes the input data and transforms it into a set of abstract representations. The decoder then takes these representations and generates the output, using both the information from the encoder and the previously generated elements in the sequence. Essentially, the encoder is responsible for understanding the input, while the decoder focuses on producing the final output.

What are the main differences between LLMs and traditional statistical language models?

ANSWER

Architecture - LLMs are based on transformers with self-attention, capturing long-range dependencies, unlike traditional models like N-grams or HMMs that struggle with this. **Scale** - LLMs have billions of parameters and train on massive datasets, enabling better generalization. Traditional models are smaller and task-specific. **Training** - LLMs undergo unsupervised pre-training and fine-tuning. Traditional models rely on supervised learning with labeled data for each task. **Context** - LLMs generate contextual embeddings, adapting to word meaning based on context. Traditional models use static embeddings.

What is a 'context window'?

ANSWER

The 'context window' in large language models (LLMs) is the span of text—measured in tokens or words—that the model can process at any given moment when generating or interpreting language. Its importance lies in its influence on the model's ability to produce coherent and contextually relevant responses. A larger context window allows the model to incorporate more surrounding information, but it also raises computational demands, creating a trade-off between performance and resource efficiency.

What is a hyperparameter?

ANSWER

A hyperparameter is a parameter that is set before the training process begins and influences how the model is trained. It controls aspects of the training process and is chosen by the developer or researcher based on prior knowledge or experimentation. Common examples include: Train-test split ratio. Learning rate in optimization algorithms (e.g., gradient descent). Choice of optimization algorithm (e.g., Adam optimizer). Choice of activation function (e.g., Sigmoid, ReLU, Tanh).

Can you explain the concept of attention mechanisms in transformer models?

ANSWER

At a high level, attention allows the model to focus on different parts of the input sequence when making predictions. Instead of treating every token equally, the model learns to 'attend' to relevant words that contribute most to the current prediction, regardless of their position. For example, in the sentence 'The dog chased the ball because it was fast,' the attention mechanism helps the model figure out that 'it' likely refers to 'the ball' based on the context.

What are Large Language Models?

ANSWER

A Large Language Model (LLM) is an AI system trained on vast amounts of text to understand, generate, and predict human-like language. It learns patterns, context, and relationships in the data to produce relevant, coherent responses. LLMs can handle a wide range of tasks, from answering questions and summarizing text to performing translations and creative writing.

What are some common challenges associated with using LLMs?

ANSWER

Using LLMs presents several common challenges:

Computational Resources - They require substantial computing power and memory for both training and deployment. **Bias and Fairness** - LLMs might learn and reproduce biases from their training data, leading to unfair outputs. **Interpretability** - It can be challenging to understand and explain how LLMs make decisions due to their complex and opaque nature. **Data Privacy** - Training on large datasets can raise concerns about data privacy and security. **Cost** - Developing, training, and deploying LLMs can be costly, which may limit their use for smaller organizations.